

# Homework 5

Name: Zheyuan Hu

NetID: zh2095

## Derivation

### Q1

The Object function is given by

$$J(w) = \lambda \|w\|^2 + \frac{1}{n} \sum_{i=1}^n \max_{y \in \mathbf{y}} [\Delta(y_i, y) + \langle w, \Psi(x_i, y) - \Psi(x_i, y_i) \rangle]$$

Since  $\Delta(y_i, y)$  is a constant of  $w$  and  $\langle w, \Psi(x_i, y) - \Psi(x_i, y_i) \rangle$  is a inner product of  $w$ ,  $\Delta(y_i, y) + \langle w, \Psi(x_i, y) - \Psi(x_i, y_i) \rangle$  is convex.

According to the Convex Optimization note 3.2.4, their pointwise maximum  $\max_{y \in \mathbf{y}} [\Delta(y_i, y) + \langle w, \Psi(x_i, y) - \Psi(x_i, y_i) \rangle]$  is also convex.

Since the sum of convex functions is also convex,

$$\sum_{i=1}^n \max_{y \in \mathbf{y}} [\Delta(y_i, y) + \langle w, \Psi(x_i, y) - \Psi(x_i, y_i) \rangle] \text{ is convex}$$

Then from note 3.1.3, we know that every norm is convex, so  $\lambda \|w\|^2$  is convex.

Therefore,  $J(w)$  is a convex function of  $w$ .

### Q2

Define

$$\hat{y}_i = \operatorname{argmax}_{y \in \mathbf{y}} [\Delta(y_i, y) + \langle w, \Psi(x_i, y) - \Psi(x_i, y_i) \rangle]$$

Then  $J(w)$  can be written as

$$J(w) = \lambda \|w\|^2 + \frac{1}{n} \sum_{i=1}^n [\Delta(y_i, \hat{y}_i) + \langle w, \Psi(x_i, \hat{y}_i) - \Psi(x_i, y_i) \rangle]$$

Since  $J(w)$  is convex and differentiable, the subgradient of  $J(w)$  is its gradient,

$$\partial J(w) = \nabla J(w) = 2\lambda w + \frac{1}{n} \sum_{i=1}^n [\Psi(x_i, \hat{y}_i) - \Psi(x_i, y_i)]$$

### Q3

An expression for the stochastic subgradient based on the point  $(x_i, y_i)$ :

$$2\lambda w + \Psi(x_i, \hat{y}_i) - \Psi(x_i, y_i)$$

### Q4

An expression for a minibatch subgradient, based on the points  $(x_i, y_i), \dots, (x_i + m - 1, y_i + m - 1)$ :

$$2\lambda w + \frac{1}{m} \sum_{j=i}^{i+m-1} [\Psi(x_j, \hat{y}_j) - \Psi(x_j, y_j)]$$

## Implementation


### Q5

```
In [6]: import numpy as np
import matplotlib.pyplot as plt
try:
    from sklearn.datasets.samples_generator import make_blobs
except:
    from sklearn.datasets import make_blobs

%matplotlib inline
```

```
In [7]: # Create the training data
np.random.seed(2)
X, y = make_blobs(n_samples=300, cluster_std=.25, centers=np.array([(-3, 1), (0, 2), (3, 1)])
plt.scatter(X[:, 0], X[:, 1], c=y, s=50)
```

Out[7]: <matplotlib.collections.PathCollection at 0x288d18c2130>



```
In [8]: from sklearn.base import BaseEstimator, ClassifierMixin, clone

class OneVsAllClassifier(BaseEstimator, ClassifierMixin):
    """
    One-vs-all classifier
    We assume that the classes will be the integers 0,...,(n_classes-1).
    We assume that the estimator provided to the class, after fitting, has a "decision"
    returns the score for the positive class.
    """
    def __init__(self, estimator, n_classes):
        """
        Constructed with the number of classes and an estimator (e.g. an
        SVM estimator from sklearn)
        @param estimator : binary base classifier used
        @param n_classes : number of classes
        """
        self.n_classes = n_classes
        self.estimators = [clone(estimator) for _ in range(n_classes)]
        self.fitted = False
```

```

def fit(self, X, y=None):
    """
    This should fit one classifier for each class.
    self.estimated[i] should be fit on class i vs rest
    @param X: array-like, shape = [n_samples,n_features], input data
    @param y: array-like, shape = [n_samples,] class labels
    @return returns self
    """

    #Your code goes here
    for i in range(self.n_classes):
        y_fit = np.array(y==i).astype(int)
        self.estimated[i].fit(X, y_fit)
    self.fitted = True
    return self

def decision_function(self, X):
    """
    Returns the score of each input for each class. Assumes
    that the given estimator also implements the decision_function method (which s
    and that fit has been called.
    @param X : array-like, shape = [n_samples, n_features] input data
    @return array-like, shape = [n_samples, n_classes]
    """

    if not self.fitted:
        raise RuntimeError("You must train classifier before predicting data.")

    if not hasattr(self.estimated[0], "decision_function"):
        raise AttributeError(
            "Base estimator doesn't have a decision_function attribute.")

    #Replace the following return statement with your code
    score = np.zeros([X.shape[0], self.n_classes])
    for i in range(self.n_classes):
        score[:,i] = self.estimated[i].decision_function(X)
    return score

def predict(self, X):
    """
    Predict the class with the highest score.
    @param X: array-like, shape = [n_samples,n_features] input data
    @returns array-like, shape = [n_samples,] the predicted classes for each input
    """

    #Replace the following return statement with your code
    score = self.decision_function(X)
    y_pred = []
    for i in range(X.shape[0]):
        y_pred.append(np.argmax(score[i,:]))
    return np.array(y_pred)

```

## Q6

I modified the regularization weight  $c=100$  to get a better separating hyperplane

```

In [9]: #Here we test the OneVsAllClassifier
from sklearn import svm
svm_estimator = svm.LinearSVC(loss='hinge', fit_intercept=False, C=100)
clf_onevsall = OneVsAllClassifier(svm_estimator, n_classes=3)
clf_onevsall.fit(X,y)

for i in range(3) :

```

```

print("Coeffs %d"%i)
print(clf_onevsall.estimators[i].coef_) #Will fail if you haven't implemented fit

# create a mesh to plot in
h = .02 # step size in the mesh
x_min, x_max = min(X[:,0])-3,max(X[:,0])+3
y_min, y_max = min(X[:,1])-3,max(X[:,1])+3
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
mesh_input = np.c_[xx.ravel(), yy.ravel()]

Z = clf_onevsall.predict(mesh_input)
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)
# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)

from sklearn import metrics
metrics.confusion_matrix(y, clf_onevsall.predict(X))

```

C:\Users\52673\anaconda3\lib\site-packages\sklearn\svm\\_base.py:985: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.  
 warnings.warn("Liblinear failed to converge, increase "

```


Coeffs 0
[[-1.05853334 -0.90294603]]
Coeffs 1
[[-0.25046466 -0.12232678]]
Coeffs 2
[[ 0.89164752 -0.82601734]]

```

```

Out[9]: array([[100,  0,  0],
               [  0, 100,  0],
               [  0,  0, 100]], dtype=int64)

```



## Q7

```

In [63]: def zeroOne(y, a) :
          ,,,
          Computes the zero-one loss.
          @param y: output class
          @param a: predicted class
          @return 1 if different, 0 if same
          ,,,
          return int(y != a)

          def featureMap(X, y, num_classes) :
            ,,,

```

```

Computes the class-sensitive features.
@param X: array-like, shape = [n_samples,n_inFeatures] or [n_inFeatures,], input features
@param y: a target class (in range 0,...,num_classes-1)
@return array-like, shape = [n_samples,n_outFeatures], the class sensitive feature map
'''

#The following line handles X being a 1d-array or a 2d-array
num_samples, num_inFeatures = (1,X.shape[0]) if len(X.shape) == 1 else (X.shape[0],X.shape[1])
#your code goes here, and replaces following return
X = np.array([X]) if len(X.shape) == 1 else X
y = np.array([y]) if type(y) != np.ndarray else y
n_outFeatures = num_classes * num_inFeatures
Psi = np.zeros([num_samples, n_outFeatures])
for i in range(num_samples):
    Psi[i, y[i]*num_inFeatures:(y[i]+1)*num_inFeatures] = X[i]
return Psi

```

## Q8

```

In [64]: def sgd(X, y, num_outFeatures, subgd, eta = 0.1, T = 10000):
'''
Runs subgradient descent, and outputs resulting parameter vector.
@param X: array-like, shape = [n_samples,n_features], input training data
@param y: array-like, shape = [n_samples,], class labels
@param num_outFeatures: number of class-sensitive features
@param subgd: function taking x,y,w and giving subgradient of objective
@param eta: learning rate for SGD
@param T: maximum number of iterations
@return: vector of weights
'''

num_samples = X.shape[0]
#your code goes here and replaces following return statement
w = np.zeros([1,num_outFeatures])
for i in range(T):
    k = np.random.randint(num_samples)
    w -= eta * subgd(X[k], y[k],w)
return w

```

## Q9

```

In [65]: class MulticlassSVM(BaseEstimator, ClassifierMixin):
'''
Implements a Multiclass SVM estimator.
'''

def __init__(self, num_outFeatures, lam=1.0, num_classes=3, Delta=zeroOne, Psi=featureMap):
'''
Creates a MulticlassSVM estimator.
@param num_outFeatures: number of class-sensitive features produced by Psi
@param lam: l2 regularization parameter
@param num_classes: number of classes (assumed numbered 0,...,num_classes-1)
@param Delta: class-sensitive loss function taking two arguments (i.e., target and predicted class)
@param Psi: class-sensitive feature map taking two arguments
'''

self.num_outFeatures = num_outFeatures
self.lam = lam
self.num_classes = num_classes
self.Delta = Delta
self.Psi = lambda X,y : Psi(X,y,num_classes)
self.fitted = False

def subgradient(self, x, y, w):
'''

```

```

Computes the subgradient at a given data point x,y
@param x: sample input
@param y: sample class
@param w: parameter vector
@return returns subgradient vector at given x,y,w
,,,

#Your code goes here and replaces the following return statement
loss = 0
for yi in range(self.num_classes):
    a = self.Delta(yi, y) + np.dot(w, (self.Psi(x, y)-self.Psi(x, yi)).T)
    if a > loss:
        loss = a
        y_hat = yi
subgrad = 2*self.lam*w + self.Psi(x, y_hat) - self.Psi(x, y)
return subgrad

def fit(self, X, y, eta=0.1, T=10000):
    ,,,

    Fits multiclass SVM
    @param X: array-like, shape = [num_samples, num_inFeatures], input data
    @param y: array-like, shape = [num_samples,], input classes
    @param eta: learning rate for SGD
    @param T: maximum number of iterations
    @return returns self
    ,,,

    self.coef_ = sgd(X, y, self.num_outFeatures, self.subgradient, eta, T)
    self.fitted = True
    return self

def decision_function(self, X):
    ,,,

    Returns the score on each input for each class. Assumes
    that fit has been called.
    @param X : array-like, shape = [n_samples, n_inFeatures]
    @return array-like, shape = [n_samples, n_classes] giving scores for each samp
    ,,,

    if not self.fitted:
        raise RuntimeError("You must train classifier before predicting data.")

    #Your code goes here and replaces following return statement
    score = np.zeros([X.shape[0], self.num_classes])
    for k in range(X.shape[0]):
        for yi in range(self.num_classes):
            Psi = self.Psi(X[k], yi)
            score[k, yi]=np.dot(self.coef_, Psi.T)
    return(score)

def predict(self, X):
    ,,,

    Predict the class with the highest score.
    @param X: array-like, shape = [n_samples, n_inFeatures], input data to predict
    @return array-like, shape = [n_samples,], class labels predicted for each data
    ,,,

    #Your code goes here and replaces following return statement
    score = self.decision_function(X)
    y_pred = []
    for i in range(X.shape[0]):
        y_pred.append(np.argmax(score[i, :]))
    return np.array(y_pred)


```

```
In [66]: #the following code tests the MulticlassSVM and sgd
#will fail if MulticlassSVM is not implemented yet
est = MulticlassSVM(6, lam=1)
est.fit(X, y, eta=0.1)
print("w:")
print(est.coef_)
Z = est.predict(mesh_input)
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)
# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)

from sklearn import metrics
metrics.confusion_matrix(y, est.predict(X))
```

```
w:
[[-1.11076154 -0.09503514  0.00667478  0.31704149  1.10408676 -0.22200634]]
```

```
Out[66]: array([[100,   0,   0],
               [  0, 100,   0],
               [  0,   0, 100]], dtype=int64)
```



```
In [ ]:
```