

Homework 6

Name: Zheyuan Hu

NetID: zh2095

Decision Tree Implementation

Load Data

```
In [1]: import matplotlib.pyplot as plt
        from itertools import product
        import numpy as np
        from collections import Counter
        from sklearn.base import BaseEstimator, RegressorMixin, ClassifierMixin
        from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, export_graphviz
        from sklearn.ensemble import GradientBoostingClassifier, GradientBoostingRegressor,
        import graphviz

        from IPython.display import Image

        %matplotlib inline
```

```
In [2]: data_train = np.loadtxt('svm-train.txt')
        data_test = np.loadtxt('svm-test.txt')
        x_train, y_train = data_train[:, 0: 2], data_train[:, 2].reshape(-1, 1)
        x_test, y_test = data_test[:, 0: 2], data_test[:, 2].reshape(-1, 1)
```

```
In [3]: # Change target to 0-1 label
        y_train_label = np.array(list(map(lambda x: 1 if x > 0 else 0, y_train))).reshape(-1, 1)
```

Q1

```
In [4]: def compute_entropy(label_array):
        """
        Calculate the entropy of given label list

        :param label_array: a numpy array of binary labels shape = (n, 1)
        :return entropy: entropy value
        """
        # Your code goes here
        K = np.unique(label_array)
        N = label_array.shape[0]
        entropy = 0
        for k in K:
            p = np.count_nonzero(label_array == k)/N
            entropy -= p*np.log(p)
        return entropy

        def compute_gini(label_array):
            """
            Calculate the gini index of label list

            :param label_array: a numpy array of labels shape = (n, 1)
```

```

: return gini: gini index value
,,,

# Your code goes here
K = np.unique(label_array)
N = label_array.shape[0]
gini = 0
for k in K:
    p = np.count_nonzero(label_array == k)/N
    gini += p*(1-p)
return gini

```

Q2

```

In [5]: class Decision_Tree(BaseEstimator):

    def __init__(self, split_loss_function, leaf_value_estimator,
                  depth=0, min_sample=5, max_depth=10):
        ,,,

        Initialize the decision tree classifier

        :param split_loss_function: method with args (X, y) returning loss
        :param leaf_value_estimator: method for estimating leaf value from array of ys
        :param depth: depth indicator, default value is 0, representing root node
        :param min_sample: an internal node can be splitted only if it contains points
        :param max_depth: restriction of tree depth.
        ,,,

        self.split_loss_function = split_loss_function
        self.leaf_value_estimator = leaf_value_estimator
        self.depth = depth
        self.min_sample = min_sample
        self.max_depth = max_depth
        self.is_leaf = False

    def fit(self, x, y):
        ,,,

        This should fit the tree classifier by setting the values self.is_leaf,
        self.split_id (the index of the feature we want ot split on, if we're splitting),
        self.split_value (the corresponding value of that feature where the split is),
        and self.value, which is the prediction value if the tree is a leaf node. If
        splitting the node, we should also init self.left and self.right to be Decision_Tree
        objects corresponding to the left and right subtrees. These subtrees should be
        the data that fall to the left and right, respectively, of self.split_value.
        This is a recursive tree building procedure.

        :param X: a numpy array of training data, shape = (n, m)
        :param y: a numpy array of labels, shape = (n, 1)

        :return self
        ,,,

        # Your code goes here
        # if reach min sample or max depth, then end up splitting
        if self.depth == self.max_depth or len(y) <= self.min_sample:
            self.is_leaf = True
            self.value = self.leaf_value_estimator(y)
            return self

        split_id, split_value = self.find_best_feature_split(x, y)
        if split_id != None and split_value != None:
            # update depth after each splitting
            self.depth += 1
            # initialize the two trees to fit the left and right data
            self.left = Decision_Tree(self.split_loss_function, self.leaf_value_estimator,

```

```

        self.right = Decision_Tree(self.split_loss_function, self.leaf_value_estimator,
                                   # split with the optimal feature and the corresponding value
                                   idx_left = np.where(x[:, split_id] <= split_value),
                                   idx_right = np.where(x[:, split_id] > split_value),
                                   x_left = x[idx_left],
                                   x_right = x[idx_right],
                                   y_left = y[idx_left],
                                   y_right = y[idx_right],
                                   # fit the tree on the left and right nodes
                                   self.left.fit(x_left, y_left),
                                   self.right.fit(x_right, y_right),
                                   self.split_id = split_id,
                                   self.split_value = split_value)
    else:
        self.is_leaf = True
        self.value = self.leaf_value_estimator(y)

    return self

def find_best_split(self, x_node, y_node, feature_id):
    """
    For feature number feature_id, returns the optimal splitting point
    for data X_node, y_node, and corresponding loss
    :param X: a numpy array of training data, shape = (n_node, 1)
    :param y: a numpy array of labels, shape = (n_node, 1)
    """
    # Your code
    n_node = x_node.shape[0]
    x = x_node[:, feature_id]
    best_loss = self.split_loss_function(y_node)
    split_value = None
    for i in range(n_node):
        idx_left = np.where(x <= x[i])
        idx_right = np.where(x > x[i])
        x_left = x_node[idx_left]
        x_right = x_node[idx_right]
        y_left = y_node[idx_left]
        y_right = y_node[idx_right]
        left_loss = self.split_loss_function(y_left)
        right_loss = self.split_loss_function(y_right)
        # loss is the weighted average of left and right loss
        loss = (len(y_left)*left_loss + len(y_right)*right_loss)/len(y_node)

        if loss < best_loss:
            best_loss = loss
            split_value = x[i]
    return split_value, best_loss

def find_best_feature_split(self, x_node, y_node):
    """
    Returns the optimal feature to split and best splitting point
    for data X_node, y_node.
    :param X: a numpy array of training data, shape = (n_node, 1)
    :param y: a numpy array of labels, shape = (n_node, 1)
    """
    # Your code
    m = x_node.shape[1]
    best_loss = self.split_loss_function(y_node)
    split_id = None
    split_value = None
    for feature_id in range(m):
        value, loss = self.find_best_split(x_node, y_node, feature_id)

```

```

        if loss < best_loss:
            split_id = feature_id
            split_value = value
            best_loss = loss
        return split_id, split_value

def predict_instance(self, instance):
    """
    Predict label by decision tree

    :param instance: a numpy array with new data, shape (1, m)

    :return whatever is returned by leaf_value_estimator for leaf containing instance
    """
    if self.is_leaf:
        return self.value
    if instance[self.split_id] <= self.split_value:
        return self.left.predict_instance(instance)
    else:
        return self.right.predict_instance(instance)

```

Q3

Decision Tree Classifier

```

In [6]: def most_common_label(y):
    """
    Find most common label
    """
    label_cnt = Counter(y.reshape(len(y)))
    label = label_cnt.most_common(1)[0][0]
    return label

```

```

In [7]: class Classification_Tree(BaseEstimator, ClassifierMixin):

    loss_function_dict = {
        'entropy': compute_entropy,
        'gini': compute_gini
    }

    def __init__(self, loss_function='entropy', min_sample=5, max_depth=10):
        """
        :param loss_function(str): loss function for splitting internal node
        """

        self.tree = Decision_Tree(self.loss_function_dict[loss_function],
                                   most_common_label,
                                   0, min_sample, max_depth)

    def fit(self, X, y=None):
        self.tree.fit(X, y)
        return self

    def predict_instance(self, instance):
        value = self.tree.predict_instance(instance)
        return value

```

```

In [8]: # Training classifiers with different depth

```

```

clf1 = Classification_Tree(max_depth=1)
clf1.fit(x_train, y_train_label)

clf2 = Classification_Tree(max_depth=2)
clf2.fit(x_train, y_train_label)

clf3 = Classification_Tree(max_depth=3)
clf3.fit(x_train, y_train_label)

clf4 = Classification_Tree(max_depth=4)
clf4.fit(x_train, y_train_label)

clf5 = Classification_Tree(max_depth=5)
clf5.fit(x_train, y_train_label)

clf6 = Classification_Tree(max_depth=20)
clf6.fit(x_train, y_train_label)

# Plotting decision regions
x_min, x_max = x_train[:, 0].min() - 1, x_train[:, 0].max() + 1
y_min, y_max = x_train[:, 1].min() - 1, x_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

f, axarr = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(10, 8))


for idx, clf, tt in zip(product([0, 1], [0, 1, 2]),
                        [clf1, clf2, clf3, clf4, clf5, clf6],
                        ['Depth = {}'.format(n) for n in range(1, 7)]):

    Z = np.array([clf.predict_instance(x) for x in np.c_[xx.ravel(), yy.ravel()]])
    Z = Z.reshape(xx.shape)

    axarr[idx[0], idx[1]].contourf(xx, yy, Z, alpha=0.4)
    axarr[idx[0], idx[1]].scatter(x_train[:, 0], x_train[:, 1], c=y_train_label, alpha=0.5)
    axarr[idx[0], idx[1]].set_title(tt)

plt.show()

```




Compare decision tree with tree model in sklearn

```
In [9]: clf = DecisionTreeClassifier(criterion='entropy', max_depth=3, min_samples_split=2)
clf.fit(x_train, y_train_label)
export_graphviz(clf, out_file='tree_classifier.dot')
```

```
In [10]: # Plotting decision regions
x_min, x_max = x_train[:, 0].min() - 1, x_train[:, 0].max() + 1
y_min, y_max = x_train[:, 1].min() - 1, x_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

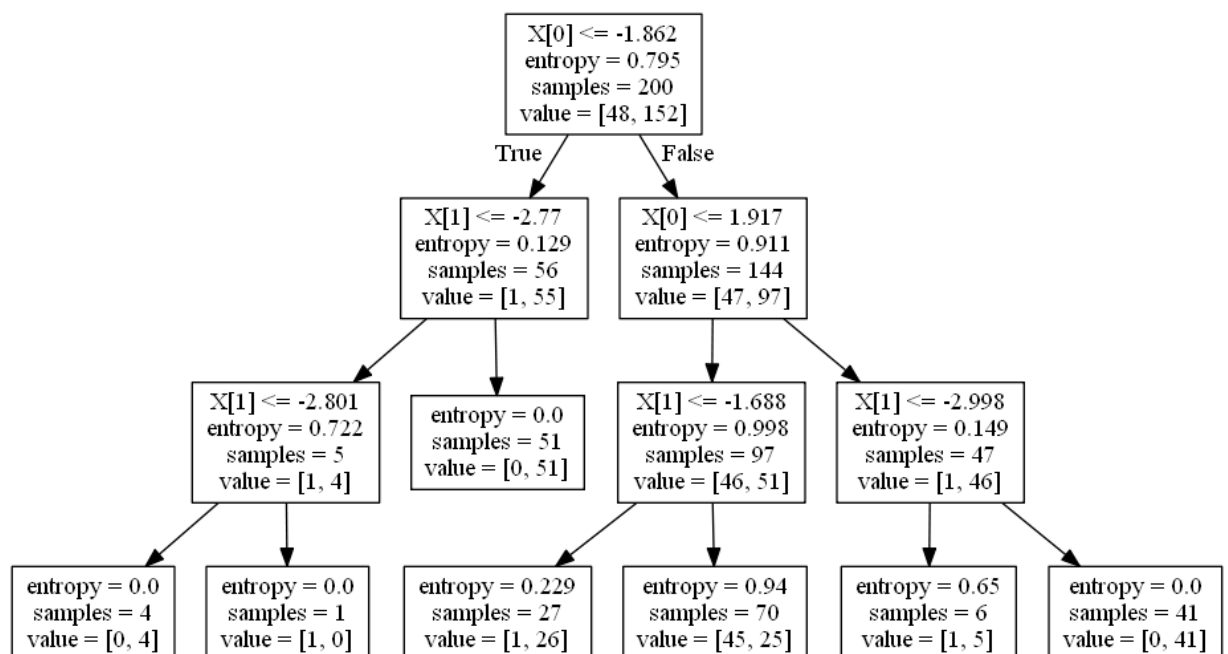
Z = np.array([clf.predict(x[np.newaxis,:]) for x in np.c_[xx.ravel(), yy.ravel()]])
Z = Z.reshape(xx.shape)
plt.figure()
plt.contourf(xx, yy, Z, alpha=0.4)
plt.scatter(x_train[:, 0], x_train[:, 1],
           c=y_train_label[:,0], alpha=0.8)
```

```
Out[10]: <matplotlib.collections.PathCollection at 0x1aa2b48d520>
```



```
In [11]: # Visualize decision tree
!dot -Tpng tree_classifier.dot -o tree_classifier.png
Image(filename='tree_classifier.png')
```

Out[11]:



Q4

Decision Tree Regressor

```
In [12]: # Regression Tree Specific Code
def mean_absolute_deviation_around_median(y):
    """
    Calculate the mean absolute deviation around the median of a given target list

    :param y: a numpy array of targets shape = (n, 1)
    :return mae
    """
    # Your code goes here
    y_med = np.median(y)
    mae = np.mean(np.abs(y - y_med))
    return mae
```

```
In [13]: class Regression_Tree():
```

```

'''
:attribute loss_function_dict: dictionary containing the loss functions used for s
:attribute estimator_dict: dictionary containing the estimation functions used in
'''

loss_function_dict = {
    'mse': np.var,
    'mae': mean_absolute_deviation_around_median
}

estimator_dict = {
    'mean': np.mean,
    'median': np.median
}

def __init__(self, loss_function='mse', estimator='mean', min_sample=5, max_depth
'''
Initialize Regression_Tree
:param loss_function(str): loss function used for splitting internal nodes
:param estimator(str): value estimator of internal node
'''

self.tree = Decision_Tree(self.loss_function_dict[loss_function],
                           self.estimator_dict[estimator],
                           0, min_sample, max_depth)

def fit(self, X, y=None):
    self.tree.fit(X, y)
    return self

def predict_instance(self, instance):
    value = self.tree.predict_instance(instance)
    return value

```

Fit regression tree to one-dimensional regression data

```

In [14]: data_krr_train = np.loadtxt('krr-train.txt')
data_krr_test = np.loadtxt('krr-test.txt')
x_krr_train, y_krr_train = data_krr_train[:,0].reshape(-1,1), data_krr_train[:,1].resh
x_krr_test, y_krr_test = data_krr_test[:,0].reshape(-1,1), data_krr_test[:,1].reshape(

# Training regression trees with different depth
clf1 = Regression_Tree(max_depth=1, min_sample=3, loss_function='mae', estimator='me
clf1.fit(x_krr_train, y_krr_train)

clf2 = Regression_Tree(max_depth=2, min_sample=3, loss_function='mae', estimator='me
clf2.fit(x_krr_train, y_krr_train)

clf3 = Regression_Tree(max_depth=3, min_sample=3, loss_function='mae', estimator='me
clf3.fit(x_krr_train, y_krr_train)

clf4 = Regression_Tree(max_depth=4, min_sample=3, loss_function='mae', estimator='me
clf4.fit(x_krr_train, y_krr_train)

clf5 = Regression_Tree(max_depth=5, min_sample=3, loss_function='mae', estimator='me
clf5.fit(x_krr_train, y_krr_train)

clf6 = Regression_Tree(max_depth=10, min_sample=3, loss_function='mae', estimator='m
clf6.fit(x_krr_train, y_krr_train)

plot_size = 0.001
x_range = np.arange(0., 1., plot_size).reshape(-1, 1)

```



```
f2, axarr2 = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(15, 10))

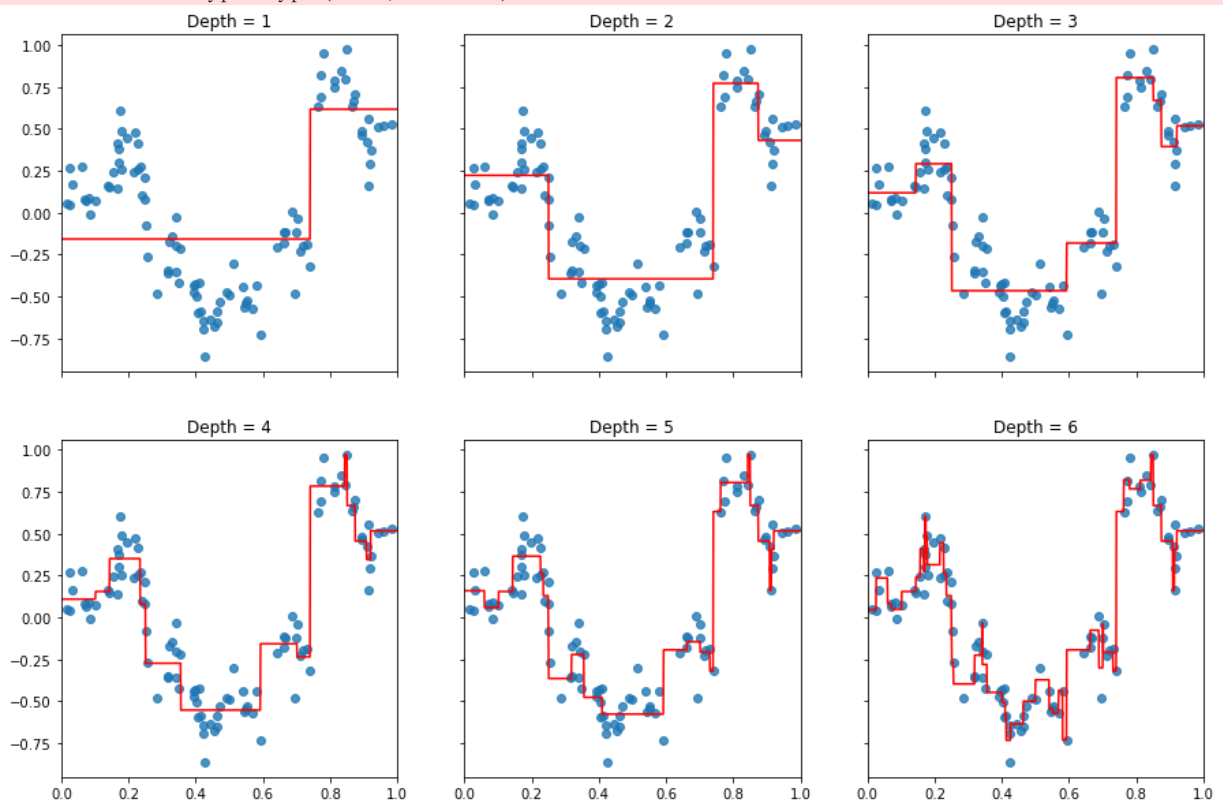
for idx, clf, tt in zip(product([0, 1], [0, 1, 2]),
                        [clf1, clf2, clf3, clf4, clf5, clf6],
                        ['Depth = {}'.format(n) for n in range(1, 7)]):

    y_range_predict = np.array([clf.predict_instance(x) for x in x_range]).reshape(-1)

    axarr2[idx[0], idx[1]].plot(x_range, y_range_predict, color='r')
    axarr2[idx[0], idx[1]].scatter(x_krr_train, y_krr_train, alpha=0.8)
    axarr2[idx[0], idx[1]].set_title(tt)
    axarr2[idx[0], idx[1]].set_xlim(0, 1)
plt.show()
```

C:\Users\52673\anaconda3\lib\site-packages\numpy\core\fromnumeric.py:3372: RuntimeWarning: Mean of empty slice.

return _methods._mean(a, axis=axis, dtype=dtype,
C:\Users\52673\anaconda3\lib\site-packages\numpy\core_methods.py:170: RuntimeWarning: invalid value encountered in double_scalars
ret = ret.dtype.type(ret / rcount)



Ensembling

Q5

```
In [15]: #Pseudo-residual function.

def pseudo_residual_L2(train_target, train_predict):
    """
    Compute the pseudo-residual based on current predicted value.
    """
    return train_target - train_predict
```

```
In [16]: class gradient_boosting():
    """
```

```

Gradient Boosting regressor class
:method fit: fitting model
,,,

def __init__(self, n_estimator, pseudo_residual_func, learning_rate=0.1, min_samp
,,,

    Initialize gradient boosting class

:param n_estimator: number of estimators (i.e. number of rounds of gradient bo
:param pseudo_residual_func: function used for computing pseudo-residual
:param learning_rate: step size of gradient descent
,,,

self.n_estimator = n_estimator
self.pseudo_residual_func = pseudo_residual_func
self.learning_rate = learning_rate
self.min_sample = min_sample
self.max_depth = max_depth

self.estimators = []

def fit(self, train_data, train_target):
,,,

    Fit gradient boosting model
,,,

    # Your code goes here
    train_predict = 0
    residuals = self.pseudo_residual_func(train_target.reshape(-1), train_predict)
    for i in range(self.n_estimator):
        hm = DecisionTreeRegressor(max_depth=self.max_depth, min_samples_leaf=self
        self.estimators.append(hm)
        self.estimators[i].fit(train_data, residuals)
        train_predict += self.learning_rate * self.estimators[i].predict(train_da
        residuals = self.pseudo_residual_func(train_target.reshape(-1), train_pre
    return self

def predict(self, test_data):
,,,

    Predict value
,,,

    # Your code goes here
    test_predict = 0
    for i in range(len(self.estimators)):
        test_predict += self.learning_rate * self.estimators[i].predict(test_data
    return test_predict

```

Q6

1-D GBM visualization - KRR data

```

In [17]: plot_size = 0.001
x_range = np.arange(0., 1., plot_size).reshape(-1, 1)

f2, axarr2 = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(15, 10))

for idx, i, tt in zip(product([0, 1], [0, 1, 2]),
                        [1, 5, 10, 20, 50, 100],
                        ['n_estimator = {}'.format(n) for n in [1, 5, 10, 20, 50, 100])

    gbm_ld = gradient_boosting(n_estimator=i, pseudo_residual_func=pseudo_residual_L2
                              max_depth=3, learning_rate=0.1)
    gbm_ld.fit(x_krr_train, y_krr_train[:, 0])

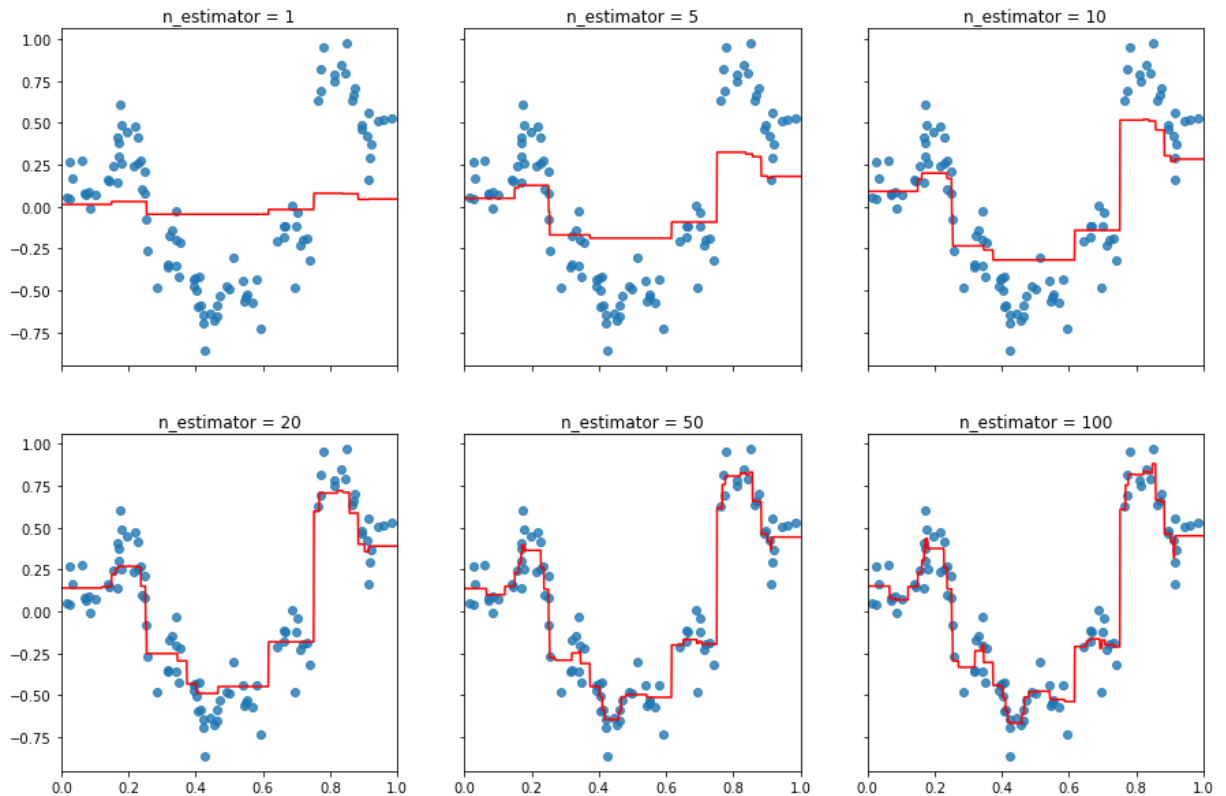
```

```

y_range_predict = gbm_ld.predict(x_range)

axarr2[idx[0], idx[1]].plot(x_range, y_range_predict, color='r')
axarr2[idx[0], idx[1]].scatter(x_krr_train, y_krr_train, alpha=0.8)
axarr2[idx[0], idx[1]].set_title(tt)
axarr2[idx[0], idx[1]].set_xlim(0, 1)

```



Classification of images with Gradient Boosting

Q7

The logistic loss:

$$\ell(y_i, f_{m-1}(x_i)) = \ln(1 + e^{-y_i f_{m-1}(x_i)})$$

The pseudo residual:

$$\begin{aligned}
 -\mathbf{g}_m &= -\left(\frac{\partial}{\partial f_{m-1}(x_j)} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i)) \right)_{j=1}^n \\
 &= -\left(\frac{\partial}{\partial f_{m-1}(x_j)} \sum_{i=1}^n \ln(1 + e^{-y_i f_{m-1}(x_i)}) \right)_{j=1}^n \\
 &= -\left(\frac{\partial}{\partial f_{m-1}(x_j)} \ln(1 + e^{-y_j f_{m-1}(x_j)}) \right)_{j=1}^n \\
 &= \left(\frac{y_j}{1 + e^{y_j f_{m-1}(x_j)}} \right)_{j=1}^n
 \end{aligned}$$

The dimension of \mathbf{g}_m :

$$\dim(\mathbf{g}_m) = n$$

Q8

$$\begin{aligned}
 h_m &= \operatorname{argmin}_{h \in \mathcal{F}} \sum_{i=1}^n ((-\mathbf{g}_m)_i - h(x_i))^2 \\
 &= \operatorname{argmin}_{h \in \mathcal{F}} \sum_{i=1}^n \left(\frac{y_i}{1 + e^{y_i f_{m-1}(x_i)}} - h(x_i) \right)^2
 \end{aligned}$$

Q9

```
In [18]: from sklearn.datasets import fetch_openml
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.utils import check_random_state

from sklearn.ensemble import GradientBoostingClassifier
```

```
In [19]: def pre_process_mnist_01():
        """
        Load the mnist datasets, selects the classes 0 and 1
        and normalize the data.
        Args: none
        Outputs:
            X_train: np.array of size (n_training_samples, n_features)
            X_test: np.array of size (n_test_samples, n_features)
            y_train: np.array of size (n_training_samples)
            y_test: np.array of size (n_test_samples)
        """
        X_mnist, y_mnist = fetch_openml('mnist_784', version=1,
                                         return_X_y=True, as_frame=False)
        indicator_01 = (y_mnist == '0') + (y_mnist == '1')
        X_mnist_01 = X_mnist[indicator_01]
        y_mnist_01 = y_mnist[indicator_01]
        X_train, X_test, y_train, y_test = train_test_split(X_mnist_01, y_mnist_01,
                                                            test_size=0.33,
                                                            shuffle=False)

        scaler = StandardScaler()
        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)

        y_test = 2 * np.array([int(y) for y in y_test]) - 1
        y_train = 2 * np.array([int(y) for y in y_train]) - 1
        return X_train, X_test, y_train, y_test
```

```
In [20]: X_train, X_test, y_train, y_test = pre_process_mnist_01()
```

Train the model

```
In [21]: estimators = [2, 5, 10, 100, 200]
train_acc_set = []
test_acc_set = []

for n_estimators in estimators:
    clf = GradientBoostingClassifier(loss='deviance', n_estimators=n_estimators, max_
```

```
clf.fit(X_train, y_train)
train_acc = clf.score(X_train, y_train)
test_acc = clf.score(X_test, y_test)
train_acc_set.append(train_acc)
test_acc_set.append(test_acc)
```

Plot the accuracy

In [22]:

```
fig, ax = plt.subplots(figsize=[10,6])
ax.set_title('Accuracy vs Num of Estimators')
ax.set_xlabel('Estimators')
ax.set_ylabel('Accuracy')
ax.plot(estimators, train_acc_set, '--x', label='train accuracy')
ax.plot(estimators, test_acc_set, '--x', label='test accuracy')
ax.legend()
plt.show()
```

